

A Priority-Based Variable Neighborhood Search with Problem-Specific Operators for Sustainable Pistachio Supply Chain Design

Raissa Gonçalves Diniz^a, André Costa Batista^{a,b}

^a*Operations Research and Complex Systems Laboratory, Universidade Federal de Minas Gerais, Brazil*

^b*Department of Electrical Engineering, Universidade Federal de Minas Gerais, Av. Antônio Carlos 6627, Belo Horizonte, 31270-010, Minas Gerais, Brazil*

Abstract

Efficient supply chain management is critical for high-value agricultural products. While mixed-integer linear programming models for the pistachio industry exist, the literature lacks efficient heuristic approaches with clear resolution methodologies. This paper presents an enhanced Variable Neighborhood Search (VNS) approach for pistachio supply chain optimization. We propose an integer-based priority encoding with a formally defined decoding algorithm. Unlike prior studies using real-valued encodings, our representation establishes a one-to-one genotype-phenotype mapping that eliminates redundancy. We develop domain-specific neighborhood structures that exploit the problem's constraint structure, improving search convergence compared to classical operators. Benchmarking against Genetic Algorithms (GA), Iterated Local Search (ILS), and an exact solver reveals that the proposed VNS outperforms other metaheuristics for medium and large-scale instances, achieving near-optimal solutions where exact methods become computationally intractable.

Keywords: Supply Chain Optimization, Metaheuristics, Variable Neighborhood Search, Sustainable Logistics

Email addresses: raissagd@ufmg.br (Raissa Gonçalves Diniz),
andre-costa@ufmg.br (André Costa Batista)

1. Introduction

The growing role of agriculture in the economic development of many countries highlights the importance of integrating this sector with other industries to stimulate overall economic growth. Agriculture not only supports food security but also acts as a major contributor to exports, generating significant income and strengthening other sectors. Among various agricultural products, pistachios stand out due to their high economic and nutritional value. Major producers such as the US, Turkey, and Iran dominate global pistachio production, remaining the leading producers worldwide in 2022 and 2023 (Statista (2023)). As a result, the strategic importance of pistachios in international trade continues to grow.

With this increasing global demand for pistachios, the need for efficient supply chain networks becomes more critical. Proper supply chain management is vital for maintaining product quality, minimizing environmental impacts, and ensuring the smooth flow of goods from production to consumers. A key challenge in pistachio production is the management of its by-products, such as hulls and shells, which, while free of heavy metals or pathogens, can degrade soil quality and contribute to environmental pollution if not properly handled (Doula et al. (2014); Bartzas and Komnitsas (2017)). Furthermore, improper disposal of these materials may lead to the growth of molds that increase the risk of aflatoxin contamination, posing serious health risks such as liver cancer (Hokmabadi (2018); Dronavalli and Kang (2019)). Therefore, eco-friendly practices for reusing agricultural waste are essential to mitigate both environmental and health hazards.

Moreover, inefficiencies in supply chain networks for food and agricultural products, including pistachios, can have significant implications for product quality and trade (Casper and Sundin (2018)). Addressing these environmental and logistical concerns is crucial to sustaining both the economic and ecological viability of the pistachio industry. In the broader context of sustainable supply chains, several studies have focused on optimizing resource use and minimizing waste through recovery and recycling practices. For example, sustainable supply chain models designed for agriculture and other industries have shown significant environmental benefits by reducing waste and reusing by-products for processes like biogas and compost production (Cheraghalipour and Roghanian (2022)). Such approaches have been successfully applied in sectors such as manufacturing and waste recovery, where efficient resource management and cost reduction have been key drivers (De-

vika et al. (2014); Zohal and Soleimani (2016)). These strategies underline the importance of sustainability in modern supply chains and the potential for agricultural industries to implement similar measures.

In related research, Rajabi-Kafshgar et al. (2023) focused on optimizing the pistachio supply chain using various algorithms, primarily population-based approaches, such as the Genetic Algorithm (GA) (Fathollahi-Fard and Hajiaghahi-Keshteli (2018); Xie et al. (2022)), among others. However, a significant limitation in their work is the lack of methodological transparency. Specifically, the authors employed a real-valued encoding scheme without detailing the decoding procedure required to transform these values into a feasible transport tree. This omission hinders reproducibility and obscures the mechanism by which constraints are handled. Furthermore, recent works have raised serious concerns about so-called metaphor-based metaheuristics, arguing that many of these algorithms offer little genuine algorithmic novelty beyond a change of terminology and that their performance is often comparable to, or worse than, that of well-established methods (Campelo and Aranha (2023); Aranha et al. (2022)). In this context, algorithms such as the Imperialist Competitive Algorithm and the Social Engineering Optimizer exemplify a broader trend where new metaphors are introduced without clear benefits over classical metaheuristics, despite adding conceptual complexity.

Addressing these gaps, this research focuses on designing an optimal supply chain network for pistachios by leveraging the Variable Neighborhood Search (VNS) algorithm (Gendreau and Potvin (2010)). While we adopt the same mixed-integer linear programming (MILP) model formulation from Rajabi-Kafshgar et al. (2023), we propose a robust overhaul of the resolution methodology. We introduce a transparent integer-based priority encoding and provide a step-by-step description of the decoding algorithm, ensuring reproducibility. Furthermore, we develop problem-specific neighborhood operators designed to exploit the particular structure of the chain. These specialized operators are contrasted with classical moves to demonstrate their efficacy in exploring the solution space.

To validate the proposed approach, we conduct a comprehensive comparative analysis against a Genetic Algorithm (GA), an Iterated Local Search (ILS), and an exact Branch-and-Cut method (using the Gurobi solver). Our findings demonstrate that the proposed VNS with specialized operators consistently outperforms the other metaheuristics in terms of solution quality. On top of that, we show that for large-scale instances, where exact methods become infeasible due to memory limitations and excessive runtime, the

proposed heuristic offers a vital trade-off, providing high-quality feasible solutions efficiently.

The remainder of this paper is organized as follows. Section 2 presents the problem definition and the mathematical formulation. Section 3 details the heuristic approach, introducing the integer encoding, the specific decoding steps, and the novel domain-specific neighborhood structures. Section 4 reports the computational experiments, offering a statistical comparison between VNS, GA, ILS, and the exact solver across varying instance sizes. Finally, Section 5 summarizes the key findings and suggests future research directions.

2. Problem Definition

The logistics network proposed by Rajabi-Kafshgar et al. (2023), illustrated in Figure 1, describes the flow of products and by-products along the pistachio supply chain. The process begins with the farmers (i), who send the harvested pistachios to the processing centers (j). At these centers, the pistachios are transformed into three main components: open-shell pistachios, raw kernels, and organic waste.

Each of these components follows specific flows. Open-shell pistachios are forwarded to pistachio factories (k), which process and distribute them to final customers (n_1). Raw kernels are sent to oil extraction centers (e), responsible for producing pistachio oil, which is directed to both final customers (n_2) and cosmetic factories (s). These, in turn, process the oil into cosmetic products that are supplied to their respective consumers (n_3). Organic waste generated in the initial processing is sent to composting centers (q), where it is transformed into compost, subsequently distributed to composting customers (m).

Although in practical terms, the produced compost could be reused by the original farmers, establishing a closed-loop system, the adopted mathematical model does not explicitly represent this feedback. Instead, compost customers are not directly associated with the initial producers, characterizing an implicit feedback cycle. This approach, already adopted in previous works (Pishvaei et al. (2010); Ahmadi and Amin (2019); Banasik et al. (2016)), preserves the structural simplicity of the model. However, the inclusion of rigorous modeling of this return loop represents a promising direction for future investigations, especially in the context of sustainable supply chains.




Figure 1: Structure of the proposed pistachio supply chain.

The mixed linear problem aims to minimize total costs (1), including opening, production, and transportation costs, subject to the following capacity and demand constraints, tolerating no shortages (see more details on the variables in Table 5 of Rajabi-Kafshgar et al. (2023)):

$$\min z_1 + z_2 + z_3 \quad (1)$$

$$\text{s.t.:} \quad \sum_{j=1}^J X_{ij} \leq Cpa_i, \quad \forall i \in I \quad (2)$$

$$\sum_{i=1}^I X_{ij} \leq Cpu_j \times U_j, \quad \forall j \in J \quad (3)$$

$$\sum_{e=1}^E Ow_{eq} + \sum_{j=1}^J Gw_{jq} \leq Cpy_q \times Y_q, \quad \forall q \in Q \quad (4)$$

$$\sum_{j=1}^J Go_{jk} \leq Cpw_k \times W_k, \quad \forall k \in K \quad (5)$$

$$\sum_{j=1}^J Gr_{je} \leq Cpr_e \times R_e, \quad \forall e \in E \quad (6)$$

$$\sum_{e=1}^E O_{en_2} \leq Cpv_s \times V_s, \quad \forall s \in S \quad (7)$$

$$\sum_{k=1}^K Go_{jk} + \sum_{e=1}^E Gr_{je} + \sum_{q=1}^Q Gw_{jq} \leq \sum_{i=1}^I X_{ij}, \quad \forall j \in J \quad (8)$$

$$\sum_{k=1}^K Go_{jk} = (1 - \beta) \sum_{i=1}^I X_{ij} \theta_1, \quad \forall j \in J \quad (9)$$

$$\sum_{e=1}^E Gr_{je} = (1 - \beta) \sum_{i=1}^I X_{ij} \theta_2, \quad \forall j \in J \quad (10)$$

$$\sum_{q=1}^Q Gw_{jq} = \sum_{i=1}^I X_{ij}\theta_3, \quad \forall j \in J \quad (11)$$

$$\theta_1 + \theta_2 + \theta_3 = 1 \quad (12)$$

$$\sum_{n_1=1}^{N_1} P_{kn_1} \leq \sum_{j=1}^J Go_{jk} \times \gamma k, \quad \forall k \in K \quad (13)$$

$$\sum_{n_2=1}^{N_2} O_{en_2} + \sum_{s=1}^S Oc_{es} \leq \sum_{j=1}^J Gr_{je} \times (1 - \lambda), \quad \forall e \in E \quad (14)$$

$$\sum_{q=1}^Q Ow_{eq} \leq \sum_{j=1}^J Gr_{je} \times \lambda, \quad \forall e \in E \quad (15)$$

$$\sum_{n_3=1}^{N_3} L_{sn_1} \leq \sum_{e=1}^E Oc_{es} \times \gamma s, \quad \forall s \in S \quad (16)$$

$$\sum_{m=1}^M D_{qm} \leq \left(\sum_{j=1}^J Gw_{jq} + \sum_{e=1}^E Ow_{eq} \right) \times \gamma q, \quad \forall q \in Q \quad (17)$$

$$\sum_{k=1}^K P_{kn_1} \geq Dp_{n_1}, \quad \forall n_1 \in N_1 \quad (18)$$

$$\sum_{e=1}^E O_{en_2} \geq Du_{n_2}, \quad \forall n_2 \in N_2 \quad (19)$$

$$\sum_{s=1}^S L_{sn_3} \geq Ds_{n_3}, \quad \forall n_3 \in N_3 \quad (20)$$

$$\sum_{q=1}^Q D_{qm} \leq Dc_m, \quad \forall m \in M \quad (21)$$

- The objective function (1) minimizes the total costs of the supply chain, including opening, transportation, and production costs. The costs are divided into three components: z_1 (22), z_2 (23), and z_3 (24), which represent facility opening costs, transportation costs, and production costs, respectively, as shown in the following equations:

$$\begin{aligned}
z_1 = & \sum_{j=1}^J Fu_j \cdot U_j + \sum_{q=1}^Q Fy_q \cdot Y_q + \sum_{k=1}^K Fw_k \cdot W_k \\
& + \sum_{e=1}^E Fr_e \cdot R_e + \sum_{s=1}^S Fv_s \cdot V_s
\end{aligned} \tag{22}$$

$$\begin{aligned}
z_2 = & \sum_{i=1}^I \sum_{j=1}^J CI_i \cdot X_{ij} + \sum_{j=1}^J \sum_{k=1}^K Cu_{1j} \cdot Go_{jk} \\
& + \sum_{j=1}^J \sum_{e=1}^E Cu_{2j} \cdot Gr_{je} + \sum_{q=1}^Q \sum_{m=1}^M Cy_q \cdot D_{qm} \\
& + \sum_{k=1}^K \sum_{n_1=1}^{N_1} Cw_k \cdot P_{kn_1} \\
& + \sum_{e=1}^E Cr_e \cdot \left(\sum_{n_2=1}^{N_2} Oc_{en_2} + \sum_{s=1}^S Oc_{es} \right) \\
& + \sum_{s=1}^S \sum_{n_3=1}^{N_3} Cv_s \cdot L_{sn_3}
\end{aligned} \tag{23}$$

$$\begin{aligned}
z_3 = & \sum_{i=1}^I \sum_{j=1}^J CX_{ij} \cdot X_{ij} + \sum_{j=1}^J \sum_{k=1}^K CK_{jk} \cdot Go_{jk} \\
& + \sum_{j=1}^J \sum_{q=1}^Q CJ_{jq} \cdot Gw_{jq} + \sum_{e=1}^E \sum_{s=1}^S CS_{es} \cdot Oc_{es} \\
& + \sum_{e=1}^E \sum_{n_2=1}^{N_2} CN_{en_2} \cdot Oc_{en_2} + \sum_{e=1}^E \sum_{q=1}^Q CQ_{eq} \cdot Ow_{eq} \\
& + \sum_{s=1}^S \sum_{n_3=1}^{N_3} Cl_{sn_3} \cdot L_{sn_3} + \sum_{k=1}^K \sum_{n_1=1}^{N_1} Cp_{kn_1} \cdot P_{kn_1} \\
& + \sum_{q=1}^Q \sum_{m=1}^M Cd_{qm} \cdot D_{qm}
\end{aligned} \tag{24}$$

- The variables R_j , Y_q , W_k , Z_e , and V_s are binary and determine whether certain elements of the model are active. For example, $R_j = 1$ indicates that facility j is operational, and $R_j = 0$ otherwise. The other variables in the model are continuous and non-negative, representing real values greater than or equal to zero. These variables model flows, quantities, and other measurable factors within the system.
- Constraints (2)-(7) address the capacity and demand limitations of each facility in the supply chain. In particular, constraint (2) limits the quantity of products transported from farmers to distribution centers. Constraints (3)-(7) ensure that the quantities sent to processing centers, composting centers, pistachio factories, oil extraction centers, and cosmetic factories do not exceed the capacities of these facilities.
- The quantity of products shipped from any facility must not exceed the orders received by that facility. Consequently, constraint (8) establishes that the quantity of products shipped from a processing center must be less than or equal to the quantity received by other facilities. Constraints (9)-(12) define the quantities of open-shell pistachios, raw kernels, and waste generated during the hulling process at each processing center. Similarly, constraints (13)-(17) govern the flows within pistachio factories, oil extraction centers, cosmetic factories, and composting centers.
- Finally, to ensure that customer demand is met, constraints (18)-(21) ensure that the quantities of products delivered to pistachio, pistachio oil, cosmetics, and compost customers satisfy or exceed their respective demands.

3. Methodology

This section details the proposed Variable Neighborhood Search (VNS) algorithm with priority-based encoding. We present the integer-based chromosome representation, the backward decoding strategy that ensures feasibility, and problem-specific neighborhood structures designed to exploit the pistachio supply chain's structural properties.

3.1. Solution Representation

Representing a solution for a complex network design problem is non-trivial. Following the principles established by Gen et al. (2006) and Rajabi-Kafshgar et al. (2023), we employ a priority-based encoding method. In this scheme, the chromosome does not directly encode the flow quantities (X_{ij}), but rather the *priority* of serving specific transportation arcs.

3.1.1. Integer Permutation vs. Real-Valued Encoding

In this work, it is used an integer permutation encoding rather than the random key (0 – 1 real value) encoding used in previous studies (Rajabi-Kafshgar et al. (2023); Fathollahi-Fard and Hajiaghahi-Keshteli (2018)).

In the real-valued approach, genes are assigned random numbers $r \in [0, 1]$, and priorities are derived from sorting these values. This representation suffers from a "many-to-one" mapping redundancy: infinite combinations of real numbers can result in the exact same priority order. For example, the vectors $v_1 = [0.1, 0.2, 0.5]$ and $v_2 = [0.8, 0.85, 0.9]$ both map to the same permutation indices (1, 2, 3). This redundancy creates "flat" regions in the search landscape where small perturbations (mutations) in the genotype produce no change in the phenotype (the solution structure).

In contrast, we employ sequential integers from 1 to Z , where Z is the total number of candidate source-depot pairs in a specific chain segment. This ordinal representation ensures a one-to-one mapping between the genotype and the solution structure. This property is crucial for the VNS algorithm, as it ensures that every neighborhood move results in a distinct solution configuration, effectively eliminating redundant search steps.

3.1.2. Chromosome Structure

The solution is represented by a composite chromosome divided into eight distinct segments (S_1 to S_8), as illustrated in Fig. 2. Each segment contains a permutation of integers representing the transport priorities for a specific part of the supply chain.

3.2. Decoding Procedure

The decoding process serves as a constructive heuristic that transforms the priority chromosome into a feasible transportation tree. As shown in Fig. 3, the decoding process begins with the flow of final products (pistachios, oil, cosmetics) from factories to consumers. This ensures that all constraints are met: consumer demand is satisfied first, propagating the requirements

S1	S2	S3	S4	S5	S6	S7	S8
K	N1	S	N3	E	N2+S	J	K
J	E	I	J	Q	M	J+E	Q

S1	S2	S3	S4	S5	S6	S7	S8
1 4 5	3 2	3 2	1	1 2 3	4 5 6	1	2 3 4
4 3	1 2 5	2 3	1	3 6	4 2 5 1	4 2 5 6 7	1 3

Figure 2: Structure of the priority-based chromosome. The solution is composed of 8 distinct integer permutations ($S_1 - S_8$), each governing the transport priority of a specific subsection in the supply chain.

upstream to factories and processing centers, ensuring that facility capacities are respected at every stage.




Figure 3: The backward decoding flow. The process prioritizes satisfying end-consumer demands first to propagate requirements upstream, ensuring feasibility throughout the chain.

Algorithm 1 details the first phase of this process, utilizing the helper function *decodingStep* (Algorithm ??) to allocate flows.

The critical logic occurs in Steps 5 and 6. Here, the algorithm reconciles diverging flows from Processing Centers, which produce multiple outputs

(Kernels for Oil vs. Open-shell for Factories) in fixed proportions. Satisfying one downstream demand often creates surplus of the other output. The decoding identifies the bottleneck demand and routes surplus to the lowest-cost destination, minimizing over-supply penalties.

Each segment of the chromosome (S_1 through S_8) is decoded using the priority-based procedure described by Gen et al. (2006), which allocates transportation flows by iterating through facilities in descending priority order. At each step, the algorithm assigns available capacity from sources to destinations with remaining demand, selecting the minimum-cost arc. This greedy strategy ensures feasible solutions while respecting all capacity and demand constraints. The complete decoding procedure is detailed in Table 2 of Gen et al. (2006).

Once the intermediate flows are established, the algorithm proceeds to the upstream and waste management stages, as detailed in Algorithm 2.

First, the exact demand for raw pistachios at the Processing Centers is calculated based on the outputs determined in Phase 1 (lines 2-5). Segment S_6 is then used to fulfill this demand from Farmers. If the aggregate capacity of farmers is insufficient, the solution is marked as infeasible.

Finally, waste management is handled in Steps 9 and 10. The waste generated at Processing and Oil Extraction centers is "pushed" to Composting Centers. Segment S_8 determines the priority of these flows, directing waste to the Composting Center with available capacity and the lowest associated cost.

3.3. The VNS Algorithm

The proposed algorithm is a hybrid metaheuristic that enhances the standard Variable Neighborhood Search (VNS) (Hansen et al. (2010)) by integrating Variable Neighborhood Descent (VND) for intensification and a Simulated Annealing (SA) acceptance criterion. This hybridization addresses the limitations of pure descent methods in the highly constrained landscape of the pistachio supply chain. The procedure is outlined in Algorithm 3.

The intensification phase employs a Variable Neighborhood Descent (VND) strategy rather than a simple local search. In this step, the algorithm iterates sequentially through the entire set of neighborhood structures. Whenever an improvement is found with any operator, the search restarts from the first operator, ensuring that the resulting solution is a local optimum with respect to all defined neighborhoods.

Algorithm 1 Chromosome decoding for the supply chain (Part 1)

```
1: procedure DECODE(data)
2:   Input: structured data containing:
3:     - Capacities ( $C$ ) and demands ( $D$ ) for all facilities
4:     - Unit transportation costs ( $T$ )
5:     - Loss rates for each process stage
6:     - Consumer demands
7:     - Chromosome segments ( $S_1, S_2, \dots, S_8$ ) representing transport priorities
8:   Output: TotalCost, Decision Variables ( $X, P, L, O, \dots$ )
9:   // Step 1: Initialize variables
10:  TotalCost  $\leftarrow 0$ 
11:  // Step 2: Decode Pistachio Factories  $\rightarrow$  Pistachio Consumers
12:   $a_1 \leftarrow$  capacities of pistachio factories
13:   $b_1 \leftarrow$  demands of pistachio consumers
14:   $c_1 \leftarrow$  transportation costs
15:  ( $cost1, P$ )  $\leftarrow$  decodingStep( $S_1, a_1, b_1, c_1$ )
16:  TotalCost  $\leftarrow$  TotalCost +  $cost1$ 
17:  // Step 3: Decode Cosmetic Factories  $\rightarrow$  Cosmetic Consumers
18:   $a_2 \leftarrow$  capacities of cosmetic factories
19:   $b_2 \leftarrow$  demands of cosmetic consumers
20:   $c_2 \leftarrow$  transportation costs
21:  ( $cost2, L$ )  $\leftarrow$  decodingStep( $S_2, a_2, b_2, c_2$ )
22:  TotalCost  $\leftarrow$  TotalCost +  $cost2$ 
23:  // Step 4: Decode Oil Extraction Centers  $\rightarrow$  Consumers
24:   $a_3 \leftarrow$  capacities of oil extraction centers
25:   $b_3 \leftarrow$  demands of oil consumers + cosmetic factories
26:   $c_3 \leftarrow$  transportation costs
27:  ( $cost3, OO_c$ )  $\leftarrow$  decodingStep( $S_3, a_3, b_3, c_3$ )
28:  Separate  $OO_c$  into  $O$  (oil consumers) and  $O_c$  (cosmetic factories)
29:  TotalCost  $\leftarrow$  TotalCost +  $cost3$ 
30:  // Step 5: Decode Processing Center  $\rightarrow$  Pistachio Factories
31:   $a_4 \leftarrow$  capacities of processing centers
32:   $b_4 \leftarrow$  adjusted demands of pistachio factories
33:   $c_4 \leftarrow$  transportation costs
34:  ( $\dots, G_o$ )  $\leftarrow$  decodingStep( $S_4, a_4, b_4, c_4$ )
35:  // Step 6: Decode Processing Center  $\rightarrow$  Oil Extraction Centers
36:   $a_5 \leftarrow$  capacities of processing centers
37:   $b_5 \leftarrow$  adjusted demands of oil extraction centers
38:   $c_5 \leftarrow$  transportation costs
39:  ( $\dots, G_r$ )  $\leftarrow$  decodingStep( $S_5, a_5, b_5, c_5$ )
40: end procedure
```

Algorithm 2 Chromosome decoding for the supply chain (Part 2)

```
1: procedure DECODE(continuation)
2:   // Step 7: Adjust demands at Processing Centers
3:   for each processing center  $j$  do
4:     Calculate required outputs based on pistachio and oil flows
5:     Update  $G_o$  and  $G_r$  to match effective capacities
6:   end for
7:   // Step 8: Decode Pistachio Farmers  $\rightarrow$  Processing Centers
8:    $a_6 \leftarrow$  capacities of pistachio farmers
9:    $b_6 \leftarrow$  updated demands at processing centers
10:   $c_6 \leftarrow$  transportation costs
11:   $(cost6, X) \leftarrow$  decodingStep( $S_6, a_6, b_6, c_6$ )
12:   $TotalCost \leftarrow TotalCost + cost6$ 
13:  // Step 9: Decode Composting Centers  $\rightarrow$  Compost Consumers
14:   $a_7 \leftarrow$  capacities of composting centers
15:   $b_7 \leftarrow$  demands of compost consumers
16:   $c_7 \leftarrow$  transportation costs
17:   $(cost7, D) \leftarrow$  decodingStep( $S_7, a_7, b_7, c_7$ )
18:   $TotalCost \leftarrow TotalCost + cost7$ 
19:  // Step 10: Assign waste flows
20:  Calculate waste from processing centers ( $J$ ) and Oil Centers ( $E$ )
21:   $a_8 \leftarrow$  available capacities of composting centers
22:   $b_8 \leftarrow$  waste demands to be disposed
23:   $c_8 \leftarrow$  transportation costs
24:   $(-, O_w) \leftarrow$  decodingStep( $S_8, a_8, b_8, c_8$ )
25:  Update  $TotalCost$ 
26:  // Step 11: Define binary opening variables
27:  Set  $U, W, R, V, Y = 1$  if corresponding flow  $> 0$ 
28:  Add fixed costs to  $TotalCost$ 
29:  return  $TotalCost$ , Decision Variables, Binary Variables
30: end procedure
```

To prevent stagnation in local optima, we replace the standard deterministic acceptance criterion with a probabilistic Simulated Annealing mechanism. A new solution generated by the VND phase is accepted as the current solution if it improves the objective function or, if it is worse, with a probability given by the Boltzmann factor $P = e^{-\Delta/T}$, where Δ is the deterioration in solution quality and T is the current temperature. This allows the search to traverse valleys in the fitness landscape.

Finally, the shaking phase applies a perturbation to the current solution to escape the basin of attraction of the local optimum. The magnitude of this perturbation is adaptive, defined in our implementation as 7% of the total decision variables in the chromosome. This value was determined through preliminary experiments testing perturbation strengths ranging from 3% to 15% on representative instances. Strengths below 5% resulted in insufficient diversification, causing the algorithm to revisit previously explored regions, while values above 10% caused excessive disruption, degrading solution quality. The 7% value provided the best balance between escaping local optima and maintaining solution structure, with a minimum threshold of 10 perturbations to ensure meaningful changes even in small instances. The specific operator used for shaking changes cyclically, resetting to the first operator upon a successful acceptance and moving to the next operator if the candidate solution is rejected.

The VND phase applies the four operators in the following fixed sequence: (1) Active Reversion, (2) Inactive-Active Swap, (3) Inactive Boost, and (4) Inactive-Active Exchange two Neighbors (IA-ETN). This ordering was designed to alternate between intensification (Active Reversion) and diversification (the remaining three operators), ensuring thorough exploration of the search space before accepting a local optimum.

3.4. Neighborhood Structures

A critical inefficiency in applying standard permutation operators to priority-based encoding is their "blindness" to the phenotypic state. In the greedy decoding process (Algorithm 1), typically only a subset of high-priority nodes is required to satisfy all demands. We term these **Active Nodes** ($\mathcal{A}$), while the remaining unused nodes are **Inactive** ($\mathcal{I}$). Standard operators often swap two Inactive nodes, changing the genotype but resulting in the exact same transportation tree and objective value.

To overcome this redundancy, our implementation explicitly tracks the set of active nodes returned by the decoder and employs four structure-aware

Algorithm 3 Hybrid VNS with SA Acceptance

Require: Initial solution x , Operators set $\mathcal{N}$, Max evaluations N_{max} , Initial Temp T_0 , Cooling rate α

```
1: Output: Best solution found  $x^*$ 
2:  $x_{curr} \leftarrow x$ ;  $x^* \leftarrow x$ 
3:  $n_{eval} \leftarrow 1$ ;  $T \leftarrow T_0$ 
4:  $k \leftarrow 0$  ▷ Operator index for shaking
5: while  $n_{eval} < N_{max}$  do
6:   // Phase 1: Intensification (VND)
7:    $x_{new} \leftarrow \text{VND\_LocalSearch}(x_{curr}, \mathcal{N})$ 
8:   Update  $n_{eval}$ 
9:   // Phase 2: Update Global Best
10:  if  $f(x_{new}) < f(x^*)$  then
11:     $x^* \leftarrow x_{new}$ 
12:  end if
13:  // Phase 3: Acceptance Criteria (SA)
14:   $\Delta \leftarrow f(x_{new}) - f(x_{curr})$ 
15:  if  $\Delta \leq 0$  or  $\text{Random}(0, 1) < e^{-\Delta/T}$  then
16:     $x_{curr} \leftarrow x_{new}$ 
17:     $k \leftarrow 0$  ▷ Reset shaking operator on acceptance
18:  else
19:     $k \leftarrow (k + 1) \pmod{|\mathcal{N}|}$  ▷ Next shaking operator
20:  end if
21:  // Phase 4: Perturbation and Cooling
22:   $T \leftarrow T \times \alpha$ 
23:   $x_{curr} \leftarrow \text{Shake}(x_{curr}, \mathcal{N}[k])$ 
24: end while
```

operators. These operators are designed to either intensify the search within the current topology or diversify it by forcing interactions between the active and inactive sets:

- **Active Reversion:** This operator functions as an intensification mechanism. It selects a segment of the chromosome containing multiple nodes from set $\mathcal{A}$ and reverses their priority values. By reordering only the suppliers that are already in use, this move fine-tunes the flow distribution to find lower-cost allocations without necessarily opening new routes or altering the network structure.
- **Inactive-Active Swap:** This is a global diversification operator. It selects one active node ($i \in \mathcal{A}$) and one inactive node ($j \in \mathcal{I}$) at random and exchanges their priorities. This forces a currently used route to be downgraded and a previously unused route to be upgraded, effectively testing global topology changes by bringing new nodes into the solution.
- **Inactive Boost:** This novel problem-specific operator addresses the "starvation" issue common in capacity-constrained problems. A low-cost supplier might be ignored simply because it sits too low in the priority list to be considered before demands are met. This operator selects a random inactive node and assigns it the highest possible priority value in the chromosome. This forces the greedy decoder to maximize flow from this specific node before considering any others, potentially revealing a completely different and superior tree structure.
- **Inactive-Active Exchange two Neighbors (IA-ETN):** This operator provides a "localized" diversification. Instead of a random global swap, it selects an inactive node and swaps it with its nearest active neighbor in terms of chromosome index distance. Unlike the global swap, which can be highly disruptive, this operator explores the "boundary" of the solution, testing if a route that is "almost" good enough (close to the active set) should replace a marginally active route.

4. Computational Experiments

This section evaluates the performance of the proposed VNS approach. The experiments are divided into two phases: first, we validate the effec-

tiveness of the proposed neighborhood structures against standard operators. Second, we benchmark the optimized VNS against a Genetic Algorithm (GA), Iterated Local Search (ILS), and the exact solver Gurobi.

4.1. Experimental Setup

The algorithms were implemented in Python and executed on a workstation with an Intel Xeon E3-1240 V2 processor (8 cores) and 32 GB of RAM running Ubuntu 24.04 LTS. The exact method was solved using Gurobi Optimizer v11.0.

The test instances were generated following the methodology proposed by Rajabi-Kafshgar et al. (2023). Each instance is characterized by equal-sized entities across all echelons (i.e., $I = J = K = E = S = Q = N_1 = N_2 = N_3 = M$), with instance sizes of 10, 30, 100, 200, 400, 800, and 1600. This results in chromosome lengths ranging from 1,700 to 27,200 variables in the priority-based encoding.

All problem parameters (facility costs, production costs, transportation costs, capacities, demands, and conversion rates) were randomly generated using uniform distributions with ranges identical to those specified in Table 7 of Rajabi-Kafshgar et al. (2023). Complete parameter specifications and the generated instances are publicly available in the project repository to ensure full reproducibility.

4.1.1. Algorithm Configurations

To ensure a fair comparison across all metaheuristics, each algorithm was allocated an identical computational budget: $N_{max} = 10n$, where n represents the total number of decision variables in the priority-based encoding. This budget scales linearly with problem size, ensuring equitable evaluation across all instances.

The specific parameter settings for each algorithm are as follows:

- **Genetic Algorithm (GA):** The GA uses a population size of $P = \lfloor \sqrt{n} \rfloor$, where n is the total number of decision variables in the priority encoding, motivated by population-sizing theory suggesting that required population grows roughly with the square root of problem size. Crossover and mutation probabilities (0.9 and 0.1, respectively) and binary tournament selection follow widely used GA settings and operators in the literature (Hussain et al. (2022)).

- **Iterated Local Search (ILS):** The ILS alternates between local search and perturbation phases. The perturbation strength scales with instance size, and the algorithm restarts when no improvement is observed after 10 consecutive non-improving iterations (Tinós et al. (2024)). Local search employs the problem-specific "Inactive-Active Swap" operator.
- **Variable Neighborhood Search (VNS):** The proposed VNS uses initial temperature $T_0 = 100$, cooling rate $\alpha = 0.995$, and perturbation strength of 7% of the chromosome length. The VND phase cycles sequentially through four problem-specific operators: (1) Active Reversion, (2) Inactive-Active Swap, (3) Inactive Boost, and (4) IA-ETN. Upon finding an improvement with any operator, the VND restarts from the first operator.
- **Exact Solver (Gurobi):** The solver was configured with a time limit of 3,600 seconds (1 hour). For instances where the optimal solution could not be found within the time limit, the best feasible solution found (Incumbent) was recorded.

All metaheuristics (VNS, GA, ILS) were executed 30 times for each instance to ensure statistical robustness.

4.2. Assessment of Neighborhood Structures

To validate the contribution of the novel problem-specific neighborhood structures, we conducted a comparative study against a baseline "Standard VNS". The Standard VNS employs only the classical permutation operators found in the literature: *Swap*, *Reversion*, *Insertion*, and *Slide* (Gendreau and Potvin (2010)).

4.2.1. Statistical Analysis

We performed a statistical validation of the results across all instance sizes. First, the Shapiro-Wilk test was applied to check for normality in the distribution of objective function values. Based on the results, we selected the appropriate significance test: the t-test for normally distributed samples and the Mann-Whitney U test for non-parametric cases. The hypothesis tested was that the Proposed VNS provides a smaller (better) objective function value than the Standard VNS ($H_1 : \mu_{Novel} < \mu_{Basic}$).

Table 1 details the p-values obtained. For small instances (10 and 30), the p-values (> 0.05) indicate that the Proposed VNS is not statistically superior. In fact, the Standard VNS performs marginally better in these simple landscapes. However, for all instances of size 100 and above, the p-values are effectively zero ($p < 0.0001$), confirming that the Proposed VNS is statistically superior with high confidence.

Table 1: Statistical tests comparing Proposed VNS vs. Standard VNS. (S-W = Shapiro-Wilk p-value).

Instance	S-W (Novel)	S-W (Basic)	Test Used	p-value	Result
10	0.0000	0.0008	Mann-Whitney	1.0000	Basic Better
30	0.0626	0.0493	Mann-Whitney	0.9811	Basic Better
100	0.5115	0.5658	t-test	0.0000	Novel Superior
200	0.0434	0.5143	Mann-Whitney	0.0000	Novel Superior
400	0.2958	0.6396	t-test	0.0000	Novel Superior
800	0.7107	0.1993	t-test	0.0000	Novel Superior
1600	0.3556	0.3454	t-test	0.0000	Novel Superior

4.2.2. Performance Magnitude

Table 2 quantifies the difference in solution quality. While standard operators are sufficient for small problems, their performance degrades as the search space expands. The proposed operators demonstrate a clear scalability advantage, with the performance gap widening in favor of the Proposed VNS as the instance size increases, reaching nearly 5% improvement in the largest instances.

Figure 4 illustrates these distributions visually. The boxplots confirm that for $N \geq 100$, the Proposed VNS not only achieves better median values but also exhibits lower variance, indicating greater robustness.

4.3. Comparative Analysis

In this section, we provide a comprehensive benchmark of the proposed VNS against the Genetic Algorithm (GA), Iterated Local Search (ILS), and the exact solver (Gurobi). Fig. ?? presents the boxplots for all methods across the tested instances. The red dashed line represents the optimal solution (or best lower bound) found by Gurobi.




Figure 4: Distribution of objective function values for Proposed VNS (Novel) vs. Standard VNS (Basic) across all instance sizes. The novel approach demonstrates superior scalability.

Table 2: Average Objective Function comparison. Gap positive indicates Proposed VNS is better.

Instance	Avg. Obj (Novel)	Avg. Obj (Basic)	Gap (%)
10	7.62×10^6	7.38×10^6	-3.05%
30	2.05×10^7	2.04×10^7	-0.44%
100	7.29×10^7	7.34×10^7	+0.71%
200	1.42×10^8	1.44×10^8	+1.45%
400	2.80×10^8	2.87×10^8	+2.52%
800	5.68×10^8	5.94×10^8	+4.58%
1600	5.67×10^8	5.95×10^8	+4.83%

4.3.1. Performance Comparison

Table 3 summarizes the average objective function values and the relative gap to the optimal solution for each method.

For the smallest instance (10), the ILS achieves the best average performance, surpassing the VNS by 2.21%. This is expected, as simple local search strategies often converge faster in very small search spaces where deep exploration is unnecessary. For Instance 30, the performance of VNS and ILS is statistically indistinguishable.

However, a clear trend emerges from Instance 100 onwards. The VNS consistently outperforms both ILS and GA, with the performance gap widening as the problem complexity increases. For the largest instances (800 and 1600), the VNS maintains a gap to the optimum of approximately 3.5%, whereas the ILS gap grows to 4.7% and the GA gap exceeds 10%. This demonstrates the superior scalability of the proposed approach.

For the two largest instances (800 and 1600), Gurobi reached the one-hour time limit without proving optimality, returning the same best incumbent solution value of 5.49×10^8 for both. These values represent upper bounds rather than proven optimal solutions.

4.3.2. Statistical Validation

To rigorously validate the performance differences, we applied the same statistical methodology described in Section 4.2. Table 4 presents the normality checks (p_{SW}) and hypothesis testing results for all pairwise comparisons between VNS and its competitors.

The hypothesis tested was $H_1 : \mu_{VNS} < \mu_{Competitor}$ (one-tailed). As shown

Table 3: Descriptive Statistics: Average Objective Function values and Gap to Reference Solution (%). Note: For instances 800 and 1600, Gurobi reached the time limit without proving optimality. Gaps are calculated relative to the best incumbent solution rather than a proven optimum.

Inst.	Reference	Average Objective Function			Gap to Reference (%)		
	Gurobi (Opt/Bound)	Proposed VNS	ILS	GA	VNS	ILS	GA
10	7.32×10^6	7.62×10^6	7.46×10^6	7.82×10^6	4.17%	1.92%	6.89%
30	2.00×10^7	2.05×10^7	2.05×10^7	2.11×10^7	2.65%	2.76%	5.39%
100	7.06×10^7	7.29×10^7	7.34×10^7	7.67×10^7	3.28%	3.97%	8.59%
200	1.38×10^8	1.42×10^8	1.44×10^8	1.49×10^8	3.40%	4.48%	8.48%
400	2.71×10^8	2.80×10^8	2.83×10^8	2.96×10^8	3.27%	4.42%	9.02%
800	5.49×10^8	5.68×10^8	5.75×10^8	6.01×10^8	3.54%	4.72%	9.53%
1600	5.49×10^8	5.68×10^8	5.74×10^8	6.04×10^8	3.50%	4.65%	10.13%

in the table, the data distribution characteristics vary with instance size, requiring different statistical treatments (parametric t-test when normality holds, non-parametric Mann-Whitney U test otherwise).

Table 4: Detailed Statistical Analysis: Normality Checks and Hypothesis Testing ($\alpha = 0.05$).

Inst.	Normality Check (p_{SW})			Hypothesis Testing ($H_1 : VNS < Competitor$)				
	VNS	ILS	GA	Comparison	Normality?	Test Used	p-value	Conclusion
10	0.0000	0.541	0.000	VNS vs ILS	No	Mann-Whitney	0.9984	ILS Superior
				VNS vs GA	No	Mann-Whitney	0.0000	VNS Superior
30	0.063	0.664	0.462	VNS vs ILS	Yes	t-test	0.2944	No Sig. Diff.
				VNS vs GA	Yes	t-test	0.0000	VNS Superior
100	0.512	0.616	0.0000	VNS vs ILS	Yes	t-test	0.0000	VNS Superior
				VNS vs GA	No	Mann-Whitney	0.0000	VNS Superior
200	0.043	0.877	0.067	VNS vs ILS	No	Mann-Whitney	0.0000	VNS Superior
				VNS vs GA	No	Mann-Whitney	0.0000	VNS Superior
400	0.296	0.295	0.625	VNS vs ILS	Yes	t-test	0.0000	VNS Superior
				VNS vs GA	Yes	t-test	0.0000	VNS Superior
800	0.711	0.882	0.439	VNS vs ILS	Yes	t-test	0.0000	VNS Superior
				VNS vs GA	Yes	t-test	0.0000	VNS Superior
1600	0.356	0.191	0.127	VNS vs ILS	Yes	t-test	0.0000	VNS Superior
				VNS vs GA	Yes	t-test	0.0000	VNS Superior

The analysis confirms two distinct performance regimes. In small instances (10-30), the simple ILS is highly efficient, and VNS provides no statistical advantage (p-value > 0.05). However, a transition occurs at Instance 100. From this point onward, the VNS is statistically superior to both ILS

and GA ($p < 0.001$). Notably, in Instance 200, the VNS distribution deviated from normality ($p = 0.043$), requiring the non-parametric Mann-Whitney test, which still confirmed the result with high confidence.

5. Conclusion

This study presented an optimization framework for the pistachio supply chain network, addressing forward products and reverse logistics of by-products. We introduced an integer-based priority encoding with a formally specified backward decoding algorithm, addressing methodological gaps in prior work.

To solve the resulting Mixed-Integer Linear Programming (MILP) problem, we proposed a hybrid Variable Neighborhood Search (VNS) algorithm. A key contribution of this work was the development of problem-specific neighborhood operators. Our experimental analysis confirmed that these structure-aware operators, which exploit the distinction between active and inactive nodes, significantly outperform standard permutation moves.

Comprehensive computational experiments benchmarked the proposed VNS against a Genetic Algorithm (GA), Iterated Local Search (ILS), and the exact solver Gurobi. The results reveal a clear performance stratification across instance sizes. For small instances ($N \leq 30$), simple local search methods such as ILS are sufficient and highly efficient, often outperforming more sophisticated approaches. However, as problem complexity increases to medium-sized instances ($N \approx 100$), the proposed VNS becomes statistically superior to both GA and ILS, demonstrating its enhanced search capability. This advantage becomes even more pronounced in large-scale instances ($N \geq 800$), where exact methods become computationally intractable due to memory and time constraints. In these challenging scenarios, the proposed VNS demonstrates remarkable scalability, consistently identifying high-quality feasible solutions that significantly outperform the evolutionary approach, establishing VNS as a promising approach for practical, large-scale supply chain optimization.

Regarding future research, two primary directions emerge. First, further variations of the problem model should be explored, such as incorporating explicit closed-loop feedback loops or stochastic parameters to better reflect agricultural uncertainties. Second, and critically, there is a need to investigate new encoding methodologies specifically designed to enable

matheuristics. While priority-based encoding is effective for pure metaheuristics, developing representations that align more closely with mathematical programming structures could unlock powerful hybrid decomposition strategies. Bridging the gap between heuristic flexibility and the precision of exact solvers remains a vital frontier for solving such complex, large-scale supply chain problems.

The source code used in this study are openly available at: https://github.com/raissagd/Pistachio_CLSC.

References

- Ahmadi, S., Amin, S.H., 2019. An integrated chance-constrained stochastic model for a mobile phone closed-loop supply chain network with supplier selection. *Journal of Cleaner Production* 226, 988–1003.
- Aranha, C., Camacho Villalón, C.L., Campelo, F., Dorigo, M., Ruiz, R., Sevaux, M., Sørensen, K., Stützle, T., 2022. Metaphor-based metaheuristics: A call for action. *Swarm Intelligence* 16, 1–6.
- Banasik, M., Stanisławska-Sachadyn, A., Sachadyn, P., 2016. A simple modification of pcr thermal profile applied to evade persisting contamination. *Journal of Applied Genetics* 57, 409–415.
- Bartzas, G., Komnitsas, K., 2017. Life cycle analysis of pistachio production in greece. *Science of the Total Environment* 595, 13–24.
- Campelo, F., Aranha, C., 2023. Lessons from the evolutionary computation bestiary. *Artificial Life* 29, 421–432.
- Casper, R., Sundin, E., 2018. Reverse logistics transportation and packaging concepts in automotive remanufacturing. *Procedia Manufacturing* 25, 154–160.
- Cheraghalipour, A., Roghanian, E., 2022. A bi-level model for a closed-loop agricultural supply chain considering biogas and compost. *Environment, Development and Sustainability* , 1–47.
- Devika, K., Jafarian, A., Nourbakhsh, V., 2014. Designing a sustainable closed-loop supply chain network based on triple bottom line approach: A comparison of metaheuristics hybridization techniques. *European Journal of Operational Research* 235, 594–615.

- Doula, M.K., Kouloumbis, P., Elaiopoulos, K., 2014. Turning wastes into valuable materials: Valorization of pistachio wastes in the agricultural sector, in: Proceedings of the SYMBIOSIS International Conference, pp. 19–21.
- Dronavalli, S., Kang, M.S., 2019. Microgametophytic selection for identifying maize with resistance to aspergillus spp. and aflatoxin. *Journal of Crop Improvement* 33, 363–371.
- Fathollahi-Fard, A.M., Hajiaghaei-Keshteli, M., 2018. A stochastic multi-objective model for a closed-loop supply chain with environmental considerations. *Applied Soft Computing* 69, 232–249.
- Gen, M., Altıparmak, F., Lin, L., 2006. A genetic algorithm for two-stage transportation problem using priority-based encoding. *OR Spectrum* 28, 337–354.
- Gendreau, M., Potvin, J.Y. (Eds.), 2010. *Handbook of Metaheuristics*. 2 ed., Springer.
- Hansen, P., Mladenović, N., Moreno Pérez, J.A., 2010. Variable neighbourhood search: Methods and applications. *Annals of Operations Research* 175, 367–407.
- Hokmabadi, H., 2018. Pistachio wastes in iran and the potential to recapture them in value chain. *Pistachio and Health Journal* 1, 1–12.
- Hussain, A., Riaz, S., Amjad, M.S., Haq, E.u., 2022. Genetic algorithm with a new round-robin based tournament selection: Statistical properties analysis. *Plos one* 17, e0274456.
- Pishvaei, M.S., Farahani, R.Z., Dullaert, W., 2010. A memetic algorithm for bi-objective integrated forward/reverse logistics network design. *Computers & Operations Research* 37, 1100–1112.
- Rajabi-Kafshgar, A., Gholian-Jouybari, F., Seyedi, I., Hajiaghaei-Keshteli, M., 2023. Utilizing hybrid metaheuristic approach to design an agricultural closed-loop supply chain network. *Expert Systems with Applications* 217, 119504.

- Statista, 2023. Leading producers of pistachios worldwide in 2022 and 2023. <https://www.statista.com/statistics/933042/global-pistachio-production-by-country/>.
- Tinós, R., Przewozniczek, M.W., Whitley, D., Chicano, F., 2024. Iterated local search with linkage learning. *ACM Transactions on Evolutionary Learning and Optimization* 4, 1–29.
- Xie, Y., Sheng, Y., Qiu, M., Gui, F., 2022. An adaptive decoding biased random key genetic algorithm for cloud workflow scheduling. *Engineering Applications of Artificial Intelligence* 112, 104879.
- Zohal, M., Soleimani, H., 2016. Developing an ant colony approach for green closed-loop supply chain network design: A case study in the gold industry. *Journal of Cleaner Production* 133, 314–337.